

Marcus Neo Jing Quan

(+65) 98538012 | marcus@neojq.com | www.neojq.com | linkedin.com/in/marcusneo | github.com/marcus-neo

Founding ML Engineer with 5+ years at an AI startup, specialising in production vision-language model fine-tuning, RLHF, RAG, and tool-calling pipelines. Built and scaled end-to-end multimodal systems — from data annotation and fine-tuning to multi-GPU serving and edge deployment — on GCP. Strong foundation in classical computer vision and 3D medical imaging for full-stack ML breadth.

EDUCATION

Imperial College London

Oct 2018 – Jun 2022

MEng Electronic & Information Engineering, First Class Honours

- Final Year Project: *Data Stream Evolution using Multiresolution Wavelet Transform* — Python-based ML research applying wavelet transforms as preprocessing for anomaly detection in data streams

Hwa Chong Institution, Singapore

Jan 2014 – Dec 2015

GCE A-Levels — AAAA

EXPERIENCE

Datature

Jul 2022 – Present

Founding Machine Learning Engineer (Team Lead)

- Led a team of 3 ML engineers across training, inference, and annotation services

Vision-Language Model Fine-Tuning Pipeline — Built from scratch, deployed on GCP

- Launched the VLM product line by building a dedicated fine-tuning service with full CI/CD, supporting [Qwen2.5-VL](#), [Qwen3-VL](#), NVILA, CosmosReason1/2, and Gemma4 at 7B–72B+ scale
- Reduced GPU memory requirements via LoRA fine-tuning with quantization; scaled multi-GPU training through tensor parallelism with OOM-resilient training loops
- Covered diverse customer use cases — VQA, Chain-of-Thought reasoning with structured evaluation metrics, freeform generation, and video training via PyAV ingestion
- Extending the pipeline with tool-calling fine-tuning support and RLHF capabilities, enabling models to learn structured function invocation and preference-based alignment
- Automated annotation labelling at scale by building the Intelliscribe caption generation microservice — a production service with JSON schema validation for structured, grounded annotation outputs

Model Hosting & Inference Service — Co-authored, deployed on Google Kubernetes Engine

- Powered production VLM and classical model inference across image, video, and 3D volumetric inputs by integrating Triton Inference Server with Numba JIT optimisation
- Shipping RAG and tool-calling capabilities into the VLM inference serving layer, enabling retrieval-augmented generation and structured function dispatch for production deployments
- Reduced annotation costs for customers by building an active learning pipeline with entropy-based metrics and bitmask annotation re-upload to the platform

Annotation Data Services — VLM & Multimodal Focus

- Built a VQA/Visual Intelligence annotation pipeline: JSONL and ViFull import/export, phrase grounding with caption/bbox support, and Neural Insights for dataset analysis
- Eliminated data migration friction via annotation import/export for 10+ formats (COCO, LabelMe, SuperAnnotate, ScaleAI, OpenAI, VoTT, Supervisely, IBM) with multipolygon, sliding window, and 3D NIfTI support; 3,800+ training datasets and 1,400+ standalone annotation exports to date
- Improved dataset quality assurance by implementing Mosaic and CutMix augmentation preview for training data visualisation

Video Annotation Services — Deployed on GKE with KEDA scale-from-zero

- Enabled frame-by-frame video annotation at scale by building video interpolation and SAM2-based tracking services with PubSub queue and multiprocessing to prevent GPU memory leaks
- Hardened production readiness with KEDA authentication, production configs, and memory-optimised caching

Model Training Service — Sole author, deployed on Google Compute Engine

- Enabled customers to train on any major framework without migration overhead by building a framework-agnostic pipeline supporting JAX, TensorFlow, and PyTorch with multi-GPU support (Data/Model Parallelism); 3,800+

models trained to date across manufacturing, medical, agriculture, sports, and military sectors

- Onboarded 16+ model architectures (YOLOv8/11/26, MaskRCNN, EfficientDet, DeepLabV3, PaliGemma, DFiNE, RF-DETR, EfficientAD) spanning classification, detection, segmentation, pose estimation, and anomaly detection
- Improved model performance on high-resolution imagery via sliding window (SAHI) training; reduced manual data preparation through class auto-balancing and multilabel classification tooling

3D Medical Imaging Pipeline — Built end-to-end using MONAI

- Unlocked a new product vertical by developing [end-to-end pipelines](#) for volumetric models (SwinUNETR, SegResNet, nnUNet, SegMamba) with NIfTI/DICOM input, RTSTRUCT output, and affine-aware evaluation; deployed into active clinical radiotherapy workflows

Model Export Service — Authored, deployed on Google Cloud Run

- Enabled deployment to [edge devices](#), mobile, and real-time applications via cross-framework export (TFLite, CoreML, TensorRT, OpenVINO, ONNX) with [Float16/Int8 quantization](#) (up to 75% model size reduction) and [model pruning](#) up to 90% of parameters; 2,700+ model artifacts exported to date

CI/CD & DevOps

- Accelerated release cycles by designing GitHub Actions CI/CD pipelines across all ML services — automated linting, testing, Docker builds, and deployment to GCP (Compute Engine, Cloud Run, GKE)
- Reduced deployment image sizes through multi-stage Docker builds, while managing CUDA runtime upgrades and Python version migrations across training, export, and inference services
- Owned and maintained PyPI and npm packages for the shared Alchemy training library and nexus-core-schema type library, ensuring version consistency across services

Datature

Apr 2021 – Jan 2022

Machine Learning Intern

- Built custom on-the-fly augmentations for the TensorFlow training pipeline and microservices for augmentation preview, dataset statistics aggregation, and hyperparameter recommendations
- Developed the full backend for Portal (neural net visualisation webapp): REST API, multi-framework model compatibility (TensorFlow, PyTorch, DarkNet), video/image predictions, and bulk inference
- Built and deployed an ML model hosting service on GCP using Docker

CORE COMPETENCIES

- **Languages:** Python, C++
- **VLM / LLM:** LoRA, Quantization, Tensor Parallelism, RLHF, RAG, Tool-Calling, HuggingFace, DeepSpeed
- **ML Frameworks:** PyTorch, TensorFlow, JAX, ONNX, MONAI, Ultralytics
- **ML Serving & Export:** Triton Inference Server, TensorRT, CoreML, TFLite, OpenVINO, Numba
- **Cloud & Infrastructure:** GCP (Compute Engine, Cloud Run, GKE, PubSub), Docker, Kubernetes, KEDA, CUDA
- **Languages (Spoken):** English (native), Mandarin (native)

ACTIVITIES & INTERESTS

- **Scuba Diving** — PADI Advanced Open Water + Nitrox certified, 50+ logged dives (Bali, Indonesia; Ko Tao, Thailand; Malapascua, Philippines; Osezaki, Japan; Tioman, Malaysia)
- **Snowboarding** — Picking it up one run at a time (Norikura Kogen, Japan; Yabuli, China)
- **Contract Bridge** — Serious hobbyist; 2/1 Game Forcing, 15-17 NT, Jacoby 2NT, Bergen Raises, RKCB 1430